

RegTrace-Agent: A Context-to-Feedback Framework for Evidence-Grounded Regulatory Response Review

Anonymous ACL submission

Abstract

Financial LLM systems are increasingly asked to judge whether an institution’s response actually resolves a prior review request, yet most financial benchmarks still evaluate static QA, extraction, or classification. We introduce RegTrace-Agent, a context-to-feedback framework for evidence-grounded response resolution. Given a reviewer request, an organizational response, and revised artifacts, the framework retrieves evidence, decomposes obligations, verifies visible gaps, converts adjudicated failures into scalar/category/natural-language feedback, and trains a guarded reviewer policy. We instantiate the framework on SEC comment-letter review and construct RegTrace-SEC, 472 public EDGAR examples with amended-filing evidence, visible-evidence labels, feedback fields, and later same-topic SEC follow-up signals. Experiments show that amended evidence is necessary, full natural-language feedback is the strongest reviewer-training signal among handwritten prompting, MIPROv2, and scalar/category GEPA variants, and a guarded verifier derived from the learned policy improves further. Later SEC follow-up comments provide external validation for the benchmark labels, suggesting that regulator-style feedback is a useful learning signal for evidence-grounded compliance reviewers.

1 Introduction

Financial NLP has rapidly moved toward LLM systems for long documents, domain-specific reasoning, and decision support. Models and benchmarks such as BloombergGPT, FinGPT, FinBen, FinanceBench, FinQA, TAT-QA, and ConvFinQA show both the promise of financial LLMs and the continuing difficulty of grounding answers in financial evidence (Wu et al., 2023; Yang et al., 2023; Xie et al., 2024; Islam et al., 2023; Chen et al., 2021; Zhu et al., 2021; Chen et al., 2022). Recent FinNLP work has similarly emphasized retrieval-augmented financial QA, SEC filing QA, financial

intelligence benchmarks, factuality judging, and evasive-answer detection (Lithgow-Serrano et al., 2025; Lai et al., 2025; Klimaszewski et al., 2025; Wu et al., 2025; Al Nuaimi et al., 2025). Most such settings, however, still present a static input and ask for a static answer, label, or report.

Many operational finance and compliance workflows have a different structure. A regulator, auditor, lender, or control function requests a correction, explanation, disclosure, or supporting artifact; the organization responds that it has revised an external document; the reviewer must decide whether visible evidence actually satisfies the request. For example, a company may reply that it revised a filing to address a requested forecast disclosure, while the amended filing only shows combined revenue projections and still omits the requested entity-level assumptions. The task is therefore not to judge whether the response sounds cooperative, but whether the repaired artifact visibly closes the original gap. We call this class of problems evidence-grounded response resolution. It is related to fact verification and document-level legal NLI (Thorne et al., 2018; Koreeda and Manning, 2021), but adds a prior reviewer request, a response that may only claim a revision, revised evidence that must be checked, and textual feedback about what remains missing.

Existing work leaves three limitations for this setting. First, benchmarks rarely ask whether an organization has repaired a documented deficiency. Second, response evaluation often has weak obligation grounding: a response can sound persuasive while amended evidence omits a named disclosure, quantification, exhibit, legal basis, or accounting analysis. Third, prompt and pipeline optimization often discards feedback structure by receiving only scalar success, although review failures are naturally described in text: what was requested, what was claimed, what evidence showed, and what gap remained.

We study these limitations through U.S. Securities and Exchange Commission (SEC) comment-letter review. SEC staff issue comments requesting expanded disclosure, clarification, revised accounting discussion, removal of unsupported statements, or additional exhibits; companies respond in writing, often with statements such as “the disclosure has been revised accordingly.” After review completion, the SEC releases staff comment letters and company responses through EDGAR. These traces often point to amended filings and sometimes include later same-topic SEC follow-up comments. Accounting research treats comment letters as economically meaningful regulatory signals (Cassell et al., 2013; Bozanic et al., 2017; Baugh and Schmardebeck, 2023; Ryans, 2021); for NLP, they provide observable request-response-evidence-feedback traces.

We introduce RegTrace-Agent, a context-to-feedback framework for evidence-grounded reviewers. It reconstructs review threads, retrieves revised-document evidence, decomposes requests into obligations, adjudicates visible evidence gaps, compiles scalar/category/full feedback, and optimizes a guarded reviewer policy. The learned reviewer is not a single hand-tuned prompt; it is the policy inside a system that gathers evidence, checks obligations, emits a guarded verdict, and uses adjudicated failures to improve the next offline reviewer.

This design makes response resolution a strong testbed for language-feedback optimization. Prompt-programming work treats LLM pipelines as declarative programs whose prompts and demonstrations can be optimized (Khatab et al., 2023); GEPA argues that natural-language feedback can expose reusable failure structure that scalar reward discards (Agrawal et al., 2026). In our setting, that structure is regulator-style: a mistake can be explained as a missing requested disclosure, unsupported revision claim, absent exhibit, or incomplete analysis.

The above limitations motivate three research questions. RQ1 asks whether public review traces can be converted into a benchmark for evidence-grounded response resolution. RQ2 asks whether full natural-language feedback is a stronger reviewer-training signal than scalar prompt search and coarse categories. RQ3 asks whether the resulting decisions generalize across threads, time, and issue categories, and whether unresolved decisions are corroborated by later regulator follow-up.

We make four contributions. First, we define evidence-grounded response resolution and introduce RegTrace-Agent, a context-to-feedback framework for evidence retrieval, obligation checking, feedback construction, and reviewer optimization. Second, we instantiate it as RegTrace-SEC, a public EDGAR benchmark with SEC comments, company responses, amended-filing evidence, labels, and feedback fields. Third, we show that full natural-language feedback is the strongest reviewer-training signal, outperforming handwritten prompting, MIPROv2, scalar/category GEPA variants, and a supervised lexical baseline; a guarded verifier distilled from the learned policy improves further. Fourth, we provide evidence ablations, independent-feedback and feedback-component ablations, prompt-evolution analysis, and systematic corroboration from real next-round SEC follow-up comments.

2 Related Work

Financial NLP and evidence-grounded evaluation. Financial NLP has moved from broad document classification toward long-document reasoning, evidence grounding, and compliance-sensitive decisions. FinQA, TAT-QA, and ConvFinQA introduced financial reasoning questions over reports, tables, and conversations (Chen et al., 2021; Zhu et al., 2021; Chen et al., 2022); FinanceBench evaluates open-book QA over company documents (Islam et al., 2023). Recent FinNLP work has similarly emphasized RAG over financial documents, accessible finance-intelligence benchmarks, factuality and causality judging, SEC filing QA, and evasive-answer detection (Lithgow-Serrano et al., 2025; Klimaszewski et al., 2025; Lai et al., 2025; Al Nuaimi et al., 2025; Wu et al., 2025). Our task differs by asking whether retrieved evidence satisfies a prior regulator request.

SEC comment letters and regulatory review. Accounting and disclosure research has studied SEC comment letters as regulatory signals, including remediation costs, disclosure effects, recurring disclosure deficiencies, and textual classification of comment-letter topics (Cassell et al., 2013; Bozanic et al., 2017; Baugh and Schmardebeck, 2023; Ryans, 2021). That literature generally treats comment letters as economic or disclosure events. We instead treat the correspondence as document-grounded supervision for LLM systems. Each review unit contains a staff request, company

187 response, amended filing evidence, and sometimes
188 a next-round SEC follow-up.

189 **Prompt optimization and language feedback.**
190 DSPy reframes LLM pipelines as declarative pro-
191 grams whose prompts and demonstrations can be
192 optimized against metrics (Khatab et al., 2023).
193 MIPROv2 and related optimizers primarily use
194 scalar feedback. A broader line studies language
195 feedback, including self-refinement, verbal rein-
196 forcement, and critique-based training (Scheurer
197 et al., 2022; Madaan et al., 2023; Shinn et al., 2023).
198 GEPA uses natural-language reflection over execu-
199 tion traces (Agrawal et al., 2026). This distinction
200 is central here because SEC response errors are
201 naturally textual: missing disclosures, unsupported
202 revision claims, absent exhibits, incomplete analy-
203 sis, or over-demand beyond the SEC request.

204 3 RegTrace-Agent Framework

205 3.1 General Review-Trace Setting

206 RegTrace-Agent applies to workflows where a re-
207 sponse claims to satisfy a prior review request by
208 modifying or supplying evidence outside the re-
209 sponse itself. Formally, each review trace contains
210 a request q , a response r , a collection of revised
211 or supporting artifacts $\mathcal{D}$, and optional later re-
212 viewer feedback u^+ . The system retrieves evidence
213 $e = R(q, r, \mathcal{D})$, extracts obligations $o_{1:m} = O(q)$,
214 and applies a verifier $V(q, r, e, o_{1:m})$ whose output
215 is either resolved or unresolved. The same trace
216 can also produce feedback for learning. A scalar
217 signal says whether the prediction is correct. A
218 category signal names a broad issue or gap type. A
219 full natural-language signal describes the request,
220 the response action, the visible evidence, and any
221 unmet requirement. This makes the framework
222 applicable beyond SEC review whenever request-
223 response-evidence traces are available.

224 3.2 SEC Instantiation

225 Each SEC example contains three test-time
226 fields: `sec_comment`, `company_response`, and
227 `retrieved_snippets` from amended filings. The
228 output is binary. `resolved` means visible amended
229 evidence covers the material request; `unresolved`
230 means a named item, quantification, exhibit, docu-
231 ment, legal analysis, accounting analysis, or other
232 material requirement remains missing. The label is
233 a visible-evidence judgment, not a legal determina-
234 tion or an inference from absence of later follow-
235 up.

Quantity	Count
Examples	472
SEC review threads	109
Companies	95
Resolved / unresolved	153 / 319
Examples with verified follow-up	236
Issue categories	7
Visible-evidence gap types	8

Table 1: Benchmark scale. Issue categories include other disclosure issues, risk factors, SPAC disclosure, revenue recognition, crypto, non-GAAP, and cybersecurity.

507 3.3 SEC Trace Construction

236

Figure 1 separates the runtime reviewer from the
feedback-learning loop. In the SEC instantiation,
we harvest UPLOAD and CORRESP materials from
EDGAR, group staff comments and company re-
sponses into review threads, and select cases where
the company claims a disclosure was added, re-
vised, removed, or filed elsewhere. We identify
amended filing candidates using filing metadata,
amendment references, page and section cues, and
response text. Candidate filings are drawn from
a window of 10 days before to 60 days after the
response date, with at most 8 candidate filings per
row. A local lexical ranker retrieves top chunks,
and an LLM reranker/adjudicator evaluates the top
3 snippets. From 589 revision-claim cases, 588
have candidate filings and snippets; 472 remain
after requiring visible evidence that directly or par-
tially addresses the SEC request. We then adjudi-
cate whether visible evidence satisfies the request
and construct labels plus feedback fields.

The retrieval component gives the agent evi-
dence to inspect. The adjudication and feedback
layer gives it something to learn from: what the
reviewer asked for, what the company claimed to
do, what the visible evidence supports, and what
remains missing. These adjudication fields are not
shown as ordinary test-time inputs. They are used
for optimization feedback, audits, and analysis.


507 3.4 Scale, Categories, and Splits

265

RegTrace-SEC contains 472 examples from 109
SEC review threads and 95 companies. The bench-
mark focuses on difficult revision-claim cases, and
the label distribution is operationally realistic: 153
examples are resolved, and 319 are unresolved.
Table 1 summarizes the corpus and split regimes.

The corpus covers seven issue categories and

RegTrace-Agent: closed-loop evidence review and feedback optimization



Test-time inputs: SEC comment, company response, retrieved amended evidence. Feedback fields are train/audit only.

Figure 1: RegTrace-Agent architecture. The runtime path parses a review trace, retrieves amended-filing evidence, checks request obligations, and emits a guarded verdict. The learning loop stores scalar, category, and natural-language feedback, lets GEPA revise the reviewer policy, and reinjects the learned policy into the obligation-review step. At test time the reviewer sees only the SEC comment, company response, and retrieved amended evidence.

273 eight visible-evidence gap types; Appendix B gives
 274 the full distribution. We use grouped, temporal, and
 275 topic-holdout regimes. The grouped split holds out
 276 entire review threads (229/104/139 train/dev/test).
 277 The temporal split trains on 2022–2023 and tests
 278 on 2024–2025 (264/64/144). The topic-holdout
 279 split tests transfer to held-out issue categories
 280 (279/68/125), including crypto and non-GAAP
 281 cases absent from training categories.

282 3.5 From Labels to Feedback

283 A standard benchmark would stop at binary labels.
 284 Our construction preserves the textual structure of
 285 each decision. For every example, adjudication
 286 records the material SEC request, the company
 287 action, evidence relevance, supporting evidence or
 288 resolution basis, visible unmet requirement, and a
 289 short rationale. This lets us compare increasingly
 290 rich optimization signals:

- 291 • scalar: correct or incorrect;
- 292 • category: correct/incorrect plus coarse issue or
 293 gap type;
- 294 • full: natural-language feedback describing the
 295 SEC request, company action, evidence coverage,

and unmet requirement.

Full feedback can state that the evidence discusses
 combined revenue projections but fails to disclose
 entity-level operating costs, product pricing, gross
 margins, and forecast limitations. This is the
 reusable decision rule that scalar metrics discard.

4 Methods

4.1 Prediction Setup

The prediction model receives the SEC comment,
 company response, and raw retrieved amended-
 filing snippets. It returns a binary label and a short
 rationale. Explicit resolved is scored as resolved.
 unresolved, partially_resolved, and other in-
 sufficient labels are normalized to unresolved. The
 rationale is required for inspection but not directly
 rewarded.

4.2 Optimization Ladder

We compare reviewer-training signals rather than
 evidence access. Baseline prompt is a hand-
 written reviewer instruction. MIPROv2 is scalar
 prompt optimization. GEPA-scalar uses reflective
 prompt optimization with correctness-only feed-
 back. GEPA-category adds coarse category feed-

back. GEPA-full uses full natural-language adjudication feedback. We also include TF-IDF, a frozen sentence-encoder scoring head, a Jev typed decision gate, and an independent-feedback GEPA ablation for reviewer-facing robustness.

All methods receive the same evidence at test time, isolating the optimization signal rather than evidence access. The key comparison is whether natural-language feedback adds value beyond prompt search itself.

4.3 Implementation Details

All main prompt-optimization runs use openai/gpt-4o-mini as the frozen prediction model with temperature 0 and 500 output tokens. GEPA uses the same model as the reflection model with temperature 0.7 and 1500 reflection tokens, seed 17, max_metric_calls=150, reflection minibatch size 4, Pareto candidate selection, and 4 evaluation threads. MIPROv2 uses the same prediction and prompt models, 8 trials, 4 prompt candidates, and at most 4 labeled and 4 bootstrapped demonstrations. The frozen encoder scorer uses all-MiniLM-L6-v2 with a logistic head over request, response, evidence, and request-evidence matching features. The Jev gate uses a choice question over resolved/unresolved. The evidence reranker/adjudicator also uses gpt-4o-mini at temperature 0. Model aliases rather than dated provider snapshots are recorded in the released artifacts.

4.4 Optimization as Feedback-Conditioned Search

We view each method as search over a prompt or prompt program p . Given an input x and label y , the frozen LLM defines $\hat{y}_p(x) = f_\theta(x; p)$, and optimization seeks a prompt with high validation reward. The methods differ in what they observe after an error. Scalar optimization observes only correctness. Category feedback adds a coarse gap type. Full feedback adds a short critique describing the request, company action, visible evidence, and unmet requirement.

The intuition is information-theoretic. Let C be the latent cause of failure, such as missing quantification or unsupported revision claim. Full feedback should help when it lowers uncertainty about C , because the optimizer can propose a more targeted semantic edit. This is not a general guarantee: a noisy or misaligned language channel can hurt. The hypothesis of RegTrace-Agent is that review traces

Input	Acc.	Ma-F1	Res-F1	Unres-F1
Response only	0.712	0.575	0.333	0.817
+ raw snippets	0.784	0.748	0.651	0.844
+ oracle summary	0.964	0.956	0.938	0.975

Table 2: Evidence ablation. Adding raw amended-filing snippets raises macro-F1 by 17.3 points. The oracle summary is an upper bound, not a deployable baseline.

are unusually well suited to language feedback because the failure causes are naturally written as reviewer critiques.

4.5 What the Optimizer Sees

During optimization, each candidate prompt is evaluated on development examples. When a prediction is incorrect, the metric returns feedback in one of three forms:

```

scalar: incorrect
category: incorrect;
gap_type=missing_specific_disclosure
full: incorrect. The SEC requested
X. The company claimed Y. The visible
evidence shows Z, but does not show
required item W.
```

Full feedback summarizes the adjudicated relationship between the SEC request, response, and retrieved evidence for optimization examples only. It does not reveal hidden evidence at test time; any advantage must come from a better learned checklist for visible-evidence review.

5 Experiments

5.1 Evidence Ablation

We first test whether amended-filing evidence matters. Table 2 evaluates a response-only condition, a raw-evidence condition, and an oracle evidence-summary condition on the grouped held-out test set.

Response-only prompting cannot verify whether a revision claim is supported. Raw snippets nearly double resolved F1, while the oracle summary shows that much of the remaining difficulty is interpreting snippets against a specific request.

5.2 Main Context-to-Feedback Results

Table 3 reports the main reviewer-training ladder on the frozen grouped split. All methods receive the same test-time context.

The majority baseline has high accuracy because unresolved examples dominate, but zero resolved F1. The handwritten prompt is conservative: it has perfect unresolved precision but misses most gaps.

Method	Acc.	Macro-F1	Res-F1	Unres-F1	Unres-P	Unres-R
Majority unresolved	0.691	0.409	0.000	0.817	0.691	1.000
Baseline prompt	0.424	0.402	0.518	0.286	1.000	0.167
MIPROv2	0.655	0.652	0.619	0.684	0.929	0.542
GEPA-scalar	0.705	0.675	0.577	0.773	0.824	0.729
GEPA-category	0.597	0.595	0.562	0.627	0.870	0.490
GEPA-full	0.784	0.756	0.674	0.839	0.867	0.812

Table 3: Frozen grouped-split leaderboard. GEPA-full gives the best macro-F1 and unresolved F1 under the same raw-evidence test-time context.

MIPROv2 improves substantially, GEPA-scalar improves recall, and GEPA-category is too coarse in this run. GEPA-full is strongest overall, improving macro-F1 by 10.4 points over MIPROv2 and 8.1 points over GEPA-scalar. This shows that full feedback is not merely adding more text; it trains the reviewer to compare requested obligations against visible evidence.

Table 4 addresses two alternate explanations and one deployment question. A TF-IDF logistic-regression classifier trained on the same raw evidence reaches only 0.523 macro-F1, so lexical supervision does not explain the GEPA-full gain. A frozen sentence encoder with a trained logistic scoring head reaches 0.667 macro-F1 without generation, and a Jev typed decision gate reaches 0.722, showing that RegTrace-Agent can support cheap triage scorers, but still trails the full feedback-trained reviewer. The guarded reviewer also improves over a monolithic prompt under both gpt-4o-mini and gpt-5.4-mini, suggesting that obligation-evidence structure matters across backbones. To test possible feedback leakage, we regenerate full textual feedback with an independent gpt-5.4-mini writer that sees only test-time fields plus the gold label, not the original adjudication rationales; GEPA with this feedback reaches 0.750 macro-F1, nearly matching the original 0.756.

We also compute five-fold out-of-fold results over review-thread groups. Mean macro-F1 is 0.698 for GEPA-full, 0.647 for MIPROv2, 0.609 for GEPA-scalar, 0.545 for GEPA-category, and 0.446 for the handwritten baseline. The fold standard deviations are nontrivial (0.151 for GEPA-full and 0.121 for MIPROv2), so we treat the single grouped split as a headline comparison and the out-of-fold analysis as stability evidence rather than as a deterministic margin claim.

5.3 Why Full Feedback Helps

The advantage of GEPA-full is not simply risk aversion. Jev reaches very high unresolved recall

Check	Condition	Macro-F1
Lexical baseline	TF-IDF, response only	0.523
Lexical baseline	TF-IDF, raw snippets	0.523
Scoring head	Frozen encoder, raw snippets	0.667
Decision gate	Jev typed choice, raw snippets	0.722
Reviewer structure	4o-mini monolithic / guarded	0.638 / 0.808
Reviewer structure	5.4-mini monolithic / guarded	0.612 / 0.732
Feedback source	GEPA-full, original feedback	0.756
Feedback source	GEPA-full, independent feedback	0.750
Diagnostic upper bound	TF-IDF, oracle evidence summary	0.858

Table 4: Robustness checks. The raw-snippet TF-IDF baseline shows that lexical supervision does not explain the main result. Frozen encoder and Jev decision scorers are cheap non-generative triage components. Multi-LLM reviewer runs show that guarded obligation-evidence structure improves over monolithic prompting under two backbones. The independent-feedback ablation shows that GEPA-full does not depend on reusing the original adjudication rationales.

(0.958) but lower resolved F1 (0.576), making it useful as a triage gate that flags likely unresolved responses. GEPA-full has lower unresolved recall but higher macro-F1, because it better balances skepticism with evidence sufficiency. It learns when a company revision claim is unsupported, but also when visible amended evidence already satisfies the SEC request.

The mechanism is visible in the feedback channel. Scalar feedback says only that a decision was wrong; category feedback names a coarse gap. Full feedback states what the SEC requested, what the company claimed, what the amended evidence showed, and what requirement remained unmet. This gives the optimizer a reusable review trace rather than a bare reward. The resulting prompt evolution shifts from generic caution to obligation-level comparison: identify material SEC request elements, inspect amended evidence for named disclosures, quantification, exhibits, and legal or accounting analysis, and distinguish partial disclosure from resolution.

The trade-off is cost and calibration. GEPA-full requires adjudicated natural-language feedback

Reviewer	Macro-F1	Res-F1	Unres-F1
Monolithic reviewer	0.645	0.614	0.676
Obligation-guided reviewer	0.619	0.613	0.624
GEPA-full optimized prompt	0.756	0.674	0.839
RegTrace guarded verifier	0.801	0.727	0.874

Table 5: Reviewer-structure ablation. The guarded verifier is a policy-distillation prototype derived from the obligation-evidence policy exposed by GEPA-full prompt evolution.

and offline optimization, so it is not the cheapest runtime gate. It can also become over-skeptical when evidence substantially satisfies a request but does not present the answer in the exact form the prompt expects. We therefore view Jev and frozen encoders as high-throughput first-pass gates, and GEPA-full as the deeper reviewer for difficult or high-impact cases.

5.4 Reviewer-Structure Ablation

The optimization ladder tests feedback richness. We also ask whether reviewer structure matters. Table 5 compares a monolithic judgment, an obligation-guided judgment, and a guarded verifier that requires every material obligation to be visibly satisfied, with GEPA-full as the optimized-prompt reference.

Simply listing obligations is not enough; the gain comes from the guard that requires visible support for each material request. We report this as a policy-distillation prototype rather than the primary optimizer comparison, since it adds human-designed structure after prompt-evolution analysis.

5.5 Feedback Component Ablation

Table 6 holds test-time input fixed and changes only optimization-time feedback.

Request-action-only feedback induces aggressive skepticism and the highest unresolved recall, but sacrifices resolved F1. Full feedback teaches both sides of the boundary: unmet requirements for unresolved cases and supporting evidence for resolved cases.

5.6 Temporal and Topic Generalization

Table 7 shows that the optimized prompt transfers beyond the grouped split. In the temporal setting, GEPA-full trains/tunes on 2022–2023 and tests on 2024–2025 material. In the topic setting, test examples include crypto and non-GAAP cases absent from training issue categories.

The topic-holdout result argues against a curated topic explanation. Per-topic breakdowns show GEPA-full improves over MIPROv2 on crypto, non-GAAP, revenue recognition, risk factors, and broad other-disclosure cases, with largest gains in revenue recognition and non-GAAP.

6 External SEC Follow-Up Corroboration

Visible-evidence labels come from an adjudication pipeline, so we test whether they align with later regulator behavior. Verified same-topic next-round SEC follow-up is not the primary label, but it is a useful corroboration signal because it can restate or narrow an unresolved gap.

We additionally audit 60 examples with an independent LLM auditor that receives the SEC comment, company response, and visible amended evidence but not the original label rationale. The audit agrees with the benchmark label in 49 of 60 cases, with mean confidence 0.897. We report this as an audit of visible-evidence label consistency, not as expert legal/accounting validation.

Among 236 examples with verified follow-up, unresolved cases are corroborated at 68.1%, while resolved cases are corroborated at only 9.2% (Table 8). Follow-up availability itself is balanced across labels (49.7% for resolved and 50.2% for unresolved; Fisher exact $p = 1.0$), so this is not merely an availability artifact.

Following reviewer-style paired analysis, we also treat later follow-up as a noisy external reference on a clean 220-example subset. The benchmark label reaches 0.756 macro-F1 against this signal, compared with 0.590 for five-fold out-of-fold GEPA-full predictions and 0.460 for the handwritten baseline. We therefore use follow-up primarily as label-validity evidence, not as proof that every GEPA-full win is separately corroborated beyond the unresolved base rate.

In one expense-disclosure case, the SEC asked for material components and trend discussion. The company claimed a revision, but retrieved evidence showed only a line item and totals; GEPA-full predicted unresolved, and the next-round SEC follow-up again requested the component breakout. SPAC transaction-background cases show the same pattern for missing negotiation details.

Feedback	Acc.	Macro-F1	Res-F1	Unres-F1	Unres-P	Unres-R
Full feedback	0.784	0.756	0.674	0.839	0.867	0.812
No evidence rationale	0.712	0.682	0.583	0.780	0.826	0.740
No unmet requirement	0.741	0.708	0.609	0.806	0.833	0.781
Request + action only	0.784	0.724	0.595	0.853	0.806	0.906

Table 6: Feedback component ablation. Request-action feedback induces aggressive unresolved detection, while full feedback gives the best macro-F1 by calibrating skepticism with evidence rationales.

Split	Method	Acc.	Macro-F1	Unres-F1	Unres-R
Time	Baseline	0.451	0.434	0.336	0.206
	MIPROv2	0.521	0.519	0.489	0.340
	GEPA-scalar	0.653	0.650	0.679	0.546
	GEPA-full	0.826	0.791	0.877	0.918
Topic	Baseline	0.440	0.440	0.426	0.274
	MIPROv2	0.576	0.565	0.635	0.484
	GEPA-scalar	0.528	0.524	0.569	0.411
	GEPA-full	0.784	0.734	0.849	0.800

Table 7: Temporal and topic-holdout generalization. GEPA-full transfers substantially better than scalar optimization.

Group	N	D/P	Rate
Unres. + f/u	160	109	68.1%
Res. + f/u	76	7	9.2%

Table 8: Real follow-up corroboration. Later same-topic SEC comments directly or partially corroborate unresolved visible-evidence labels far more often than resolved labels.

Method	Claim insuff.	Evidence required	Request match	Missing items
GEPA-full	2/2	2/2	2/2	2/2
GEPA-scalar	2/2	2/2	0/2	0/2
GEPA-category	2/2	2/2	0/2	0/2
Req.+action only	2/2	2/2	2/2	2/2

Table 9: Rule presence in optimized prompt candidates. Full feedback more consistently induces named-request matching and missing-item reasoning.

7 Prompt Evolution and Case Studies

7.1 Prompt Evolution

Prompt evolution shows that GEPA-full does not merely make the prompt longer. It turns adjudication feedback into a compact regulatory reading policy. The seed prompt is a single sentence; the final GEPA-full run saves optimized candidates of 349 and 397 words (Appendix A).

Scalar and category feedback learn general caution: do not trust a revision claim unless evidence supports it. Full feedback more consistently learns named-request matching and missing-item reasoning (Table 9), shifting from generic skepticism to obligation-level comparison: SEC request → company action → visible evidence → remaining gap.

7.2 Method-Disagreement Cases

The strongest GEPA-full wins involve concrete request-evidence mismatches. In a financial-projections case, the SEC requested individual-company forecast information and material assumptions; the amended evidence showed only combined revenue projections. GEPA-full marked it unresolved because the requested assumption set was still absent.

In a crypto terminology case, the amended evidence defined stablecoins but did not explain the SEC’s requested contrast with unstable cryptocurrency. Baseline, MIPROv2, GEPA-scalar, and GEPA-category predicted resolved; GEPA-full predicted unresolved. In a PIPE financing case, MIPROv2 treated the presence of price terms as sufficient; GEPA-full required equivalent price-per-share disclosure for each agreement, a gap later echoed in SEC follow-up.

GEPA-full also has informative errors. In one ownership-percentage case, retrieved evidence contained redemption-scenario calculations and the benchmark label was resolved. GEPA-full predicted unresolved because it wanted a more explicit cover-page presentation. This reveals the main calibration risk: full-feedback prompts learn useful skepticism but can over-require exact presentation. Appendix C lists additional cases.

8 Discussion

The main result is not simply that one optimizer beats another. RegTrace-Agent turns a recurring

ring compliance pattern into an offline feedback-learning system: collect the trace, retrieve the revised artifact, check obligations against visible evidence, issue a guarded verdict, and recycle adjudicated failures as feedback. The learned reviewer can be used directly, distilled into a guarded verifier, or converted into lightweight scoring heads for high-throughput triage. The output is a traceable unresolved-gap rationale, not only a binary label.

Language feedback helps because it preserves decision structure. A scalar signal says only that a verdict failed; full feedback says which named request, company action, amended snippet, and remaining obligation were mismatched. The optimized prompt therefore becomes an interpretable policy: extract material requests, distrust unsupported revision claims, inspect evidence for named items, avoid over-requiring, and distinguish partial disclosure from resolution.

SEC correspondence is the public setting where we can observe request, response, revised filing, and later reviewer behavior, but the trace structure is broader. Audit remediation, regulatory reporting, lending-document stipulations, data-quality exceptions, and internal compliance monitoring also contain request-response-evidence loops. The transferable claim is that such traces can support retrieval, obligation checking, adjudication feedback, optimized reviewing, and hard-case mining; the limitation is that high-recall language-feedback policies can become over-skeptical.

9 Conclusion

We introduced RegTrace-Agent, a context-to-feedback framework for evidence-grounded response review, and instantiated it as RegTrace-SEC using public SEC traces. Across grouped, temporal, and topic-holdout evaluations, full natural-language feedback yields the strongest reviewer. Evidence ablations, prompt evolution, and SEC follow-up corroboration suggest that financial NLP benchmarks can move from static labels toward feedback-rich review systems that preserve the reasons behind expert decisions.

Limitations and Data Statement

This study is conditioned on visible retrieved evidence, so missed snippets can affect labels. The benchmark is LLM-adjudicated, with independent audit and external follow-up corroboration but without a full expert legal/accounting audit. A full

independent gpt-5.4-mini audit over all 472 examples, using only test-time fields, reaches 0.731 agreement accuracy and 0.715 macro-F1 against the released labels. Same-topic SEC follow-up is a noisy behavioral signal, not ground truth. We frame labels as visible-evidence judgments rather than legal determinations. An anonymized release will be provided at <https://anonymous.open.science/r/regtrace-sec-XXXX/>; it includes benchmark JSONL, splits, accession identifiers, retrieval metadata, snippets, labels, feedback fields, prompts, outputs, and analysis scripts.

666 References

- 667 Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu,
668 Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav
669 Singhvi, Herumb Shandilya, Michael J. Ryan, Meng
670 Jiang, Christopher Potts, Koushik Sen, Alexandros G.
671 Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and
672 Omar Khattab. 2026. Gepa: Reflective prompt evolu-
673 tion can outperform reinforcement learning. *ICLR*
674 2026.
- 675 Khaled Al Nuaimi, Gautier Marti, Alexis Marchal, and
676 Andreas Henschel. 2025. Detecting evasive answers
677 in financial q&a: A psychological discourse taxon-
678 omy and lightweight baselines. In *Proceedings of the*
679 *Joint Workshop on Financial Technology and Natural*
680 *Language Processing*. Association for Computational
681 Linguistics.
- 682 Matthew D. Baugh and Roy Schmardebeck. 2023. Au-
683 ditor style and common disclosure issues: Evidence
684 from sec comment letters. *The Accounting Review*.
- 685 Zahn Bozanic, J. Richard Dietrich, and Bret A. Johnson.
686 2017. Sec comment letters and firm disclosure. *Jour-*
687 *nal of Accounting and Public Policy*, 36(5):337–357.
- 688 Cory A. Cassell, Lauren M. Dreher, and Linda A. Myers.
689 2013. Reviewing the sec’s review process: 10-k
690 comment letters and the cost of remediation. *The*
691 *Accounting Review*, 88(6):1875–1908.
- 692 Zhiyu Chen, Wenhu Chen, Charese Smiley, Sameena
693 Shah, Iana Borova, Dylan Langdon, Reema Moussa,
694 Matt Beane, Ting-Hao Huang, Bryan Routledge, and
695 William Yang Wang. 2021. Finqa: A dataset of nu-
696 merical reasoning over financial data. In *Proceed-*
697 *ings of the 2021 Conference on Empirical Methods in*
698 *Natural Language Processing*. Association for Com-
699 putational Linguistics.
- 700 Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma,
701 Sameena Shah, and William Yang Wang. 2022. Con-
702 vfinqa: Exploring the chain of numerical reasoning
703 in conversational finance question answering. In *Pro-*
704 *ceedings of the 2022 Conference on Empirical Meth-*
705 *ods in Natural Language Processing*. Association for
706 Computational Linguistics.
- 707 Pranab Islam, Anand Kannappan, Douwe Kiela, Re-
708 becca Qian, Nino Scherrer, and Bertie Vidgen. 2023.
709 [Financebench: A new benchmark for financial ques-](#)
710 [tion answering](#). *Preprint*, arXiv:2311.11944.
- 711 Omar Khattab, Arnav Singhvi, Paridhi Maheshwari,
712 Zhiyuan Zhang, Keshav Santhanam, Saiful Haq,
713 Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam,
714 Heather Miller, Matei Zaharia, and Christopher
715 Potts. 2023. [Dspy: Compiling declarative language](#)
716 [model calls into self-improving pipelines](#). *Preprint*,
717 arXiv:2310.03714.
- 718 Mateusz Klimaszewski, Pinzhen Chen, Liane Guillou,
719 Ioannis Papaioannou, Barry Haddow, and Alexandra
720 Birch. 2025. Avenibench: Accessible and versatile
721 evaluation of finance intelligence. In *Proceedings*
of the Joint Workshop on Financial Technology and
Natural Language Processing. Association for Com-
putational Linguistics.
- Yuta Koreeda and Christopher Manning. 2021. [Con-](#)
[tractNLI: A dataset for document-level natural lan-](#)
[guage inference for contracts](#). In *Findings of the*
Association for Computational Linguistics: EMNLP
2021, pages 1907–1919, Punta Cana, Dominican Re-
public. Association for Computational Linguistics.
- Viet Lai, Michael Krumdick, Charles Lovering, Varshini
Reddy, Craig Schmidt, and Chris Tanner. 2025. Sec-
qa: A systematic evaluation corpus for financial qa.
In *Proceedings of the Joint Workshop on Financial*
Technology and Natural Language Processing. Asso-
ciation for Computational Linguistics.
- Oscar Lithgow-Serrano, David Kletz, Vani Kanjiran-
gat, David Adametz, Marzio Lunghi, Claudio Bone-
sana, Matilde Tristany-Farinha, Yuntao Li, Detlef
Repplinger, Marco Pierbattista, Stefania Stan, and
Oleg Szehr. 2025. Assessing rag system capabilities
on financial documents. In *Proceedings of the Joint*
Workshop on Financial Technology and Natural Lan-
guage Processing. Association for Computational
Linguistics.
- Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler
Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon,
Nouha Dziri, Shrimai Prabhumoye, Yiming Yang,
Shashank Gupta, Bodhisattwa Prasad Majumder,
Katherine Hermann, Sean Welleck, Amir Yazdan-
bakhsh, and Peter Clark. 2023. Self-refine: Iterative
refinement with self-feedback. In *Advances in Neu-*
ral Information Processing Systems.
- James P. Ryans. 2021. Textual classification of sec com-
ment letters. *Review of Accounting Studies*, 26:37–
80.
- Jeremy Scheurer, Jon Ander Campos, Jun Shern Chan,
Angelica Chen, Kyunghyun Cho, and Ethan Perez.
2022. [Training language models with language feed-](#)
[back](#). *Preprint*, arXiv:2204.14146.
- Noah Shinn, Federico Cassano, Edward Berman, Ash-
win Gopinath, Karthik Narasimhan, and Shunyu Yao.
2023. Reflexion: Language agents with verbal rein-
forcement learning. In *Advances in Neural Informa-*
tion Processing Systems.
- James Thorne, Andreas Vlachos, Christos
Christodoulopoulos, and Arpit Mittal. 2018.
[FEVER: a large-scale dataset for fact extraction](#)
[and VERification](#). In *Proceedings of the 2018*
Conference of the North American Chapter of
the Association for Computational Linguistics: Human
Language Technologies, Volume 1 (Long
Papers), pages 809–819, New Orleans, Louisiana.
Association for Computational Linguistics.
- Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski,
Mark Dredze, Sebastian Gehrmann, Prabhanjan Kam-
badur, David Rosenberg, and Gideon Mann. 2023.
[Bloomberggpt: A large language model for finance](#).
Preprint, arXiv:2303.17564.

780	Yilin Wu, Han Yuan, Li Zhang, and Zheng Ma. 2025.	
781	Natural language inference as a judge: Detecting	
782	factuality and causality issues in language model self-	
783	reasoning for financial analysis. In <i>Proceedings of</i>	
784	<i>the Joint Workshop on Financial Technology and</i>	
785	<i>Natural Language Processing</i> . Association for Com-	
786	putational Linguistics.	
787	Qianqian Xie, Weiguang Han, Zhengyu Chen, Ruoyu	
788	Xiang, Xiao Zhang, Yueru He, Mengxi Xiao, Dong	
789	Li, Yongfu Dai, Duanyu Feng, Yijing Xu, and 1 oth-	
790	ers. 2024. <i>Finben: A holistic financial benchmark for</i>	
791	<i>large language models</i> . <i>Preprint</i> , arXiv:2402.12659.	
792	Hongyang Yang, Xiao-Yang Liu, and Christina Dan	
793	Wang. 2023. <i>Fingpt: Open-source financial large</i>	
794	<i>language models</i> . <i>Preprint</i> , arXiv:2306.06031.	
795	Fengbin Zhu, Wenqiang Lei, Youcheng Huang, Chao	
796	Wang, Shuo Zhang, Jiancheng Lv, Fuli Feng, and	
797	Tat-Seng Chua. 2021. Tat-qa: A question answering	
798	benchmark on a hybrid of tabular and textual content	
799	in finance. In <i>Proceedings of the 59th Annual Meet-</i>	
800	<i>ing of the Association for Computational Linguistics</i> .	
801	Association for Computational Linguistics.	
802	A Full Optimized Prompts	
803	The main paper quotes short excerpts from opti-	
804	mized prompts. Below are the full GEPA-full op-	
805	timized candidates exposed by the final grouped-	
806	split run. The seed prompt was: “Classify whether	
807	a company resolved an SEC comment using the	
808	visible evidence.”	
809	A.1 GEPA-full Optimized Candidate 1	
810	You are tasked with classifying whether a com-	
811	pany has resolved a comment from the SEC based	
812	on the provided evidence. Follow these instruc-	
813	tions carefully:	
814	1. Input Format: You will receive three types of	
815	inputs: <code>sec_comment</code> , <code>company_response</code> , and	
816	<code>amended_evidence</code> .	
817	2. Classification Criteria: Determine if the com-	
818	pany has fully resolved the SEC comment by ana-	
819	lyzing the company response in conjunction with	
820	the amended evidence. If the response satisfies	
821	the SEC’s request and the evidence supports this	
822	claim, classify the comment as resolved. If the re-	
823	sponse does not meet the SEC’s request or if there	
824	are still outstanding issues, classify the comment	
825	as unresolved.	
826	3. Evaluation Process: Carefully compare the	
827	SEC request against the company’s response and	
828	the amended evidence. Look for specific items or	
829	information requested by the SEC, quantification	
830	or detailed analysis that meets SEC requirements,	
831	references to regulatory guidance or previous de-	
832	cisions, and clear explanations of how the com-	
833	pany adjusted disclosures or practices. Take note	
834	of visible unmet requirements, partial fixes, or	
835	lack of supporting documentation in the amended	
836	evidence.	
	4. Output Format: Include a reasoning section,	837
	a label indicating resolved or unresolved, and a	838
	concise reason.	839
	5. Feedback Incorporation: Be vigilant about	840
	claims made by the company; ensure that	841
	amended evidence corroborates these claims.	842
	Identify instances where the response may par-	843
	tially address SEC concerns but still falls short of	844
	complete compliance.	845
	A.2 GEPA-full Optimized Candidate 2	846
	You are tasked with classifying whether a com-	847
	pany has resolved a comment issued by the SEC	848
	based on provided evidence. Follow these specific	849
	instructions for the task:	850
	1. Input Format: You will receive <code>sec_comment</code> ,	851
	<code>company_response</code> , and <code>amended_evidence</code> .	852
	The SEC comment is the original request for clar-	853
	ification or additional information. The company	854
	response indicates what changes were made or	855
	explanations provided. The amended evidence	856
	contains snippets from the amended filing or other	857
	documents that support the response.	858
	2. Classification Criteria: Determine whether the	859
	company’s response and amended evidence ade-	860
	quately address the SEC’s concerns as outlined	861
	in the original comment. A response is resolved	862
	if it fully satisfies the requirements of the SEC	863
	comment, including clear explanations, specific	864
	details, and quantifications when necessary. A	865
	response is unresolved if there are still gaps or	866
	missing elements in the company’s explanation or	867
	if amended evidence does not address all parts of	868
	the SEC comment.	869
	3. Evaluation Process: Carefully analyze the SEC	870
	comment and identify the specific details or expla-	871
	nations requested. Review the company response	872
	to see if the company acknowledged the comment	873
	and provided a meaningful response. Examine	874
	amended-evidence snippets to confirm that they	875
	support the company’s claims and cover all re-	876
	quested aspects of the SEC comment. Look for	877
	direct references in the amended evidence that	878
	address the SEC’s specific requests, such as own-	879
	ership structures, accounting adjustments, or en-	880
	forcement risks.	881
	4. Labeling: If the company has fully addressed	882
	the SEC’s concerns with adequate evidence, label	883
	the response resolved. If there are omissions or	884
	inadequacies in addressing SEC requests, label it	885
	unresolved.	886
	5. Justification: Provide clear and concise rea-	887
	soning that cites specific elements from the SEC	888
	comment, company response, and amended evi-	889
	dence. Highlight specific unmet requirements or	890
	gaps in the evidence that contribute to an unre-	891
	solved classification.	892
	6. Feedback Incorporation: Pay attention to feed-	893
	back from previous examples to identify common	894
	pitfalls, such as failing to compare the SEC re-	895
	quest against the evidence adequately or missing	896
	key elements in the response.	897
	B Additional Dataset Details	898
	The most common issue categories are other disclo-	899
	sure issues (183), risk factors (91), SPAC disclosure	900

901 (78), revenue recognition (75), crypto (28), non-
902 GAAP (16), and cybersecurity (1). The most com-
903 mon gap types are resolved visible evidence (153),
904 missing specific disclosure (140), missing quantifi-
905 cation (59), resolved-but-regulator-requested-more-
906 detail (38), missing exhibit or document (29), in-
907 complete accounting analysis (25), incomplete le-
908 gal or regulatory analysis (25), and other visible
909 evidence gap (3).

910 **C Additional Qualitative Cases**

911 The qualitative appendix used during analysis in-
912 cludes full SEC comments, company responses,
913 amended evidence snippets, method predictions,
914 GEPA-full rationales, and later SEC follow-up text
915 for selected cases. These include stablecoin termi-
916 nology, financial forecast assumptions, PIPE price-
917 per-share disclosure, Item 404 related-party disclo-
918 sure, expense breakout and trend discussion, and
919 SPAC transaction-background disclosure.